BIOMETRIC EVALUATION COMMON FRAMEWORK

PROGRAMMER'S GUIDE
VERSION 0.1

WAYNE SALAMON
GREGORY FIUMARA

IMAGE GROUP
INFORMATION ACCESS DIVISION
INFORMATION TECHNOLOGY LABORATORY

NIST

National Institute of
Standards and Technology
U.S. Department of Commerce

JANUARY 16, 2014

Contents

Chapter 1

Introduction

This document describes the Biometric Evaluation Framework (BECommon) and application programming interfaces (API) used to support the evaluation of biometric software within the NIST Image Group [13].

1.1   Rationale

When evaluating software in a "black box" fashion many aspects of program execution must be addressed, such as non-returning function calls, I/O errors, and other resource requirements. In addition, solutions to common problems should be portable across operating systems.

An evaluation consists of the testing of vendor-supplied software that implements certain biometric algorithms, such as fingerprint matching or face recognition. The NIST Image Group defines a test process and API for each evaluation. Vendors implement the API in their software, which is delivered to NIST as a software library, where common test driver is used to call the vendor library to perform the biometric operation. In order to support the common functionality used across all evaluations, such as logging, file input/output, etc., a common framework is used.

Even though the Biometric Evaluation Framework was written to support biometric software evaluations, much of the framework can be used for any general purpose programs where data storage and system interaction are needed. One goal of the BECommon is to reduce the low-level error processing (particularly with input and output) done directly by applications. The Biometric Evaluation Framework provides several abstractions that are useful to applications so they can focus on the task at hand.

This document describes the BECommon in two sections: Chapters containing descriptions of each package as well as code examples, and reference sections containing auto-generated API documentation.

The BECommon is a work-in-progress, and future development will occur in areas where the need arises for the testing programs of the NIST Image Group.

Chapter 2

Overview

The Biometric Evaluation Framework (BECommon) is a set of C++[15] classes, error codes, and design patterns used to create a common environment to provide logging, data management, error handling, and other functionality that is needed for many applications used in the testing of biometric software. The goals of the framework include:

- Reduce the amount of I/O error handling implemented by applications.

- Provide standard interfaces for data management and logging;

- Remove the need for applications to handle low-level events from the operating system (signals, etc.);

- Provide services for timing the execution of code blocks;

- Allow applications to constrain the amount of processing time used by a block of code.

The experience of the NIST Image Group when running many software evaluations has led to the need of a common code for dealing with recurring software issues. One issue is the large amounts of data consumed, and created, by the software under test. Input data sets are typically biometric images, while output sets contain derived information. Both sets of data often contain millions of items, and storing each item as a file creates a tremendous burden on the file system. The *IO* package provides a solution to managing large amounts of records in a portable, efficient manner, as well as facilities for logging and maintaining runtime settings.

BECommon is divided into several packages, each providing a set of related functionality, such as error handling and timing operations. The packages are an informal concept, mapped to formal C++ name spaces, e.g. `IO` and `Time`. A namespace contains classes, constants, and non-class functions that relate to concepts grouped in the namespace. All classes within BECommon belong to the top-level `BiometricEvaluation` namespace.

Biometric image data is often supplied in a compressed format (e.g. WSQ, JPEG) and must be converted to a "raw" format. The `Image` package contains classes to represent compressed image data as an object, storing the image size and other attributes, in addition to the raw image.

Memory management issues are addressed by the `Memory` package. The use of classes and templates in this package can relieve applications of the need to directly manage memory for dynamically sized arrays, or call functions that are already provided to allocate and free C library objects.

While a program is running, it is often necessary to record certain statistics about the process, such as memory and processor usage. The `Process` package provides methods to obtain this information, as well as the capability to log to a file periodically, in an asynchronous manner.

In addition to its own statistics, a program may need to query some information about the environment under which it is running. The `System` package provides a count of CPUs, memory size, other system characteristics that an application can use to tailor its behavior.

Many aspects of software performance evaluation involve the use of timers. The `Time` package provides for the calculation of a time interval in a manner that is consistent across platforms, abstracting the underlying operating system's timing facility. Also, included is a "watchdog" facility, providing a solution to the problem of non-returning function calls. By using a watchdog timer, an application can abort a call to a function that doesn't return in the required interval.

The `Text` package provides a set of utility functions for operating on strings. The `digest` functions are of interest to those applications that must mask any information contained in a string before passing that information to another function. For example, often the biometric image file (or record) names contain information about the image, such as the finger position.

Error propagation and handling are addressed by the `Error` package. A set of exception objects are defined within this package, allowing for communication of error conditions out of the framework to the application, along with an explanatory string. Signal handling is related to error propagation in that when a process receives a signal, often it is due to software bug. Divide by zero, for example. The Error package provides for simple handling of the signal by the process.

Many packages in BECommon deal with biometric data record formats, including ANSI/NIST [3] records. In order to provide a general interface to several formats, BECommon represents the biometric data as derived from a source. For example, the `Finger` package contains classes that represent all information about a finger, including the source image and derived minutiae points. The `View` package combines the notions of a source image and derived information together into a single abstraction.

BECommon is designed to be used in a modular fashion, and it is possible to compile many packages independently. However, several packages do make use of other packages in the framework, and therefore, are less flexible in their reuse. However, BECommon is designed to reduce the intra-framework dependencies.

A set of test programs is included with the framework. These programs not only exercise the functions provided by the packages, but also can be used as example programs on how to use framework.

The chapters that follow this overview describe each package in detail, along with some code examples. The final set of chapters of this document contain the application programming interfaces for the types, methods, and classes that make up BECommon. However, the framework is under development, and other packages, classes, etc. will be added over time to address the needs of the NIST Image Group.

Chapter 3

Framework

The `Framework` package is used to retrieve information about the Biometric Evaluation Framework itself. Version numbers, the compiler used, and other information can be queried by applications. Versioning information is recorded in the BECommon `Makefile` and populated in the function implentation at compile-time.

Listing 3.1: Using the `Framework` API

```cpp
#include <iostream>

#include <be_framework.h>

using namespace BiometricEvaluation;
using namespace std;

int
main(
    int argc,
    char* argv[])
{
        cout << "Framework Version: ";
        cout << Framework::getMajorVersion() << "." <<
            Framework::getMinorVersion() << endl;
        /* "Framework Version: 0.4" */

        cout << "Compiler Used: ";
        cout << Framework::getCompiler() << " v" <<
            Framework::getCompilerVersion() << endl;
        /* "Compiler Used: clang v3.0.0" */

        cout << "Date/Time Compiled: ";
        cout << Framework::getCompileDate() << " " <<
            Framework::getCompileTime() << endl;
        /* "Date/Time Compiled: Jan 24 2012 12:16:01" */

        return (EXIT_SUCCESS);
}
```

Chapter 4

Memory

To assist applications with memory management, the `Memory` package provides classes to wrap C memory allocations, and other dynamically-sized objects.

4.1 AutoBuffer

The Biometric Evaluation Framework is designed to interoperate with existing C code that has its own memory management techniques, e.g. NIST Biometric Image Software [12]. In these cases, functions exist to allocate and free blocks of memory, and these calls must be made by the applications which use those libraries. To assist BECommon clients that use these existing libraries, the `AutoBuffer` class wraps the C memory management functions, guaranteeing the release of C objects when the AutoBuffer goes out of scope.

The `AutoBuffer` constructor takes three function pointers as parameters: one for C object construction, one for destruction, and a third, optional, function for copying the C object. If the latter is passed a NULL, the `AutoBuffer` and the underlying C object cannot be copied, and an exception will be thrown.

Listing 4.1 shows the use of `AutoBuffer` to wrap the memory allocation routines that are part of the NIST Biometric Image Software ANSI/NIST library.

Listing 4.1: Using the `AutoBuffer`

```
#include <be_memory_autobuffer.h>
#include <iostream>
extern "C" {
  #include <an2k.h>
}

int
main(int argc, char* argv[]) {


    /*
     * alloc_ANSI_NIST(), free_ANSI_NIST(), and copy_ANSI_NIST()
     * are functions in the NBIS AN2K library.
     */
    Memory::AutoBuffer<ANSI_NIST> an2k =
        Memory::AutoBuffer<ANSI_NIST>(&alloc_ANSI_NIST,
            &free_ANSI_NIST, &copy_ANSI_NIST);
    if (read_ANSI_NIST(fp, an2k) != 0) {
            cerr << "Could not read AN2K file." << endl;
            return (EXIT_FAILURE);
```

```
21        }
22
23        for (int i = 1; i < an2k->num_records; i++) {
24            // process the ANSI/NIST record ...
25        }
26 }
```

4.2 AutoArray

At its simplest level, `AutoArray` is a C-style array with numerous convenience methods, such as being able to query the number of elements. C++ iterators can be used over the contents of the array. The array can be resized without the need to create a new object. C++ operator overloading allows `AutoArray` objects to be passed to C-style functions that expect pointers to `AutoArray`'s template type.

`AutoArray` is used extensively in BECommon to help eliminate mistakes when manually allocating memory. The `AutoArray` constructor will allocate needed memory using `new` and the destructor will `delete` it. This ensures that any allocated memory will be appropriately freed when the `AutoArray` goes out of scope. Copy constructors and methods as well as the assignment operator all correctly manage memory so the client does not have to. Several objects in BECommon `return AutoArray` objects to assist clients in proper memory management.

A common use of `AutoArray` is to deal with records sequenced from a `RecordStore`. Listing 4.2 demonstrates this. Notice the omission of memory management statements – they are completely unnecessary.

Listing 4.2: Using **`AutoArray`** s with **`RecordStore`** s

```
1  #include <be_io_dbrecstore.h>
2  #include <be_memory_autoarray.h>
3
4  #include <iostream>
5
6  using namespace BiometricEvaluation;
7
8  int
9  main(
10     int argc,
11     char *argv[])
12 {
13         IO::DBRecordStore rs("db_recstore", ".", IO::READONLY);
14
15         uint64_t value_size = 0;
16         string key("");
17         Memory::AutoArray<uint8_t> value;
18         for (bool stop = false; stop == false; ) {
19                 try {
20                         // Non-destructively resize the AutoArray to hold
21                         // the next record.
22                         value.resize(rs.sequence(key, NULL));
23
24                         // Read the record into the AutoArray (treats the
25                         // AutoArray as a pointer).
26                         rs.read(key, value);
27
28                         // Do something with value.
29                         std::cout << "Key " << key << " has a value of " <<
30                             value.size() << " bytes" << std::endl;
```

```
31                      } catch (Error::ObjectDoesNotExist) {
32                              stop = true;
33                      }
34              }
35
36          return (0);
37 }
```

AutoArray is adapted from "c_array " [15, 496].

4.3   IndexedBuffer

Many applications have a need to read items from a data record and take action based on the value of the item read. For example, when reading a biometric data record, the number of finger minutiae points in the record is indicated by a value in the record header. Furthermore, the record format may be of a different endianess than the application's host platform.

The IndexedBuffer class is used to access data from a buffer in fixed-size amounts in sequence. Objects of this class maintain an index into the buffer as internal state and reads out of the buffer, when using certain methods, adjust the index. In addition, standard subscript access can be done on on the buffer (reads and writes) without affecting the index. The basic element type is an unsigned eight-bit value. The IndexedBuffer object can be created to either manage the buffer memory directly, or to "wrap" an existing buffer.

Methods to retrieve elements from the buffer are defined in the class's interface. These functions are used to retrieve 8/16/32/64-bit values while moving the internal index. Several functions are also provided to take into account the endianess of the underlying data.

Listing 4.3 shows how an application can read a data record in big-endian format.

Listing 4.3: Using the **IndexedBuffer**

```
1  #include <be_memory_autoarray.h>
2  #include <be_memory_indexedbuffer.h>
3
4  int
5  main(int argc, char* argv[]) {
6
7          uint64_t size = IO::Utility::getFileSize("BiometricRecord");
8          FILE *fp = std::fopen("BiometricRecord", "rb");
9          Memory::IndexedBuffer iBuf(size);
10         fread(iBuf, 1, size, fp);
11         fclose(fp);
12         Memory::IndexedBuffer iBuf(recordData, recordData.size());
13
14         uint32_t lval;
15         uint16_t sval;
16
17         /*
18          * Record is big-endian:
19          * ---------------------------------
20          * | NAME | LENGTH | ID |    ...    |
21          * ---------------------------------
22          *     4       4      2
23          */
24
25         /* Read a 4-byte C string */
26         lval = iBuf.scanU32Val();         /* Format ID */
27         char *cptr = (char *)&lval;
```

```
28          string s(cptr);
29
30          /* Read a 4-byte length */
31          lval = iBuf.scanBeU32Val()
32
33          /* Read a 2-byte ID */
34          sval = iBuf.scanBeU16Val();
35 }
```

Chapter 5

Error Handling

Within the Biometric Evaluation Framework, Error handling has two aspects: One for communicating error conditions out of the framework and back to applications; the other for handling error signals from the environment and operating system. Classes and other code to implement error processing are described in this chapter.

5.1 Biometric Evaluation Exceptions

The Biometric Evaluation Framework contains a set of classes used to report errors to applications. Objects of these class types are thrown and contain descriptive information as to the nature of the error. Applications must handle the errors in a manner that makes sense for the application.

Applications should catch objects of the type specified in the API for the class being called. The type of object caught indicates the nature of the error that occurred, while the string stored within that object provides more information on the error.

Listing 6.2 on page 17 shows an example of exception handling when using the logging classes described in Section 6.3 on page 16.

5.2 Signal Handling

When the application process executes in a POSIX environment, signals to the process can be generated by the operating system. In many cases, if the signal is not handled by the process, execution terminates. Because the Biometric Evaluation Framework was designed to used with software libraries for which no source code is available, changes to the code in these libraries cannot be made, and any faults in that code cannot be fixed. A common problem is that a function in the "black box" library dereferences a bad pointer, resulting in a segmentation violation signal being sent by the operating system.

To prevent termination of the application process, signal handling must be installed. The Biometric Evaluation Framework provides a class, `SignalManager`, to simplify the installation of a signal handler in order to allow the program to continue running. For example, when extracting a fingerprint minutia template from an image, often the library call will fault on a certain image. By using the `SignalManager`, the application can log that fault, and continue on to the next image.

Signal handling in a POSIX environment covers the bare essentials, and one of two actions is usually taken. The signal can be handled and processing continues at the location the signal was generated. The second action is that, in addition to signal handling, the process continues from a different location. It is the second action that is implemented by the `SignalManager` class. The rationale for this type of signal handling is so the call to the faulting function can be aborted, but the caller can detect that the signal was handled and take action, usually by logging the fault.

By default, the `SignalManager` class installs a handler for the `SIGSEGV` and `SIGBUS` signals. However, other signals can be handled as desired.

One restriction on the use of `SignalManager` is that the POSIX calls for signal management (`signal(3)`, `sigaction(2)`, etc.) cannot be invoked inside of the signal handler block.

The example in Listing 5.1 shows application use of the `SignalManager` class.

Listing 5.1: Using the **SignalManger**

```
1  #include <be_error_signal_manager.h>
2  using namespace BiometricEvaluation;
3
4  int main(int argc, char *argv[])
5  {
6          Error::SignalManager *sigmgr = new Error::SignalManager();
7
8          BEGIN_SIGNAL_BLOCK(sigmgr, sigblock1);
9          // code that may result in signal generation
10         END_SIGNAL_BLOCK(asigmgr, sigblock1);
11         if (sigmgr->sigHandled()) {
12                 // log the event, etc.
13         }
14 }
```

Within the `SignalManager` header file, two macros are defined: `BEGIN_SIGNAL_BLOCK()` and `END_SIGNAL_BLOCK()`, each taking the `SignalManager` object and label as parameters. The label must be unique for each signal block. These macros insert the jump buffer into the code, which is the location where the signal handler will jump to after handling the signal. The use of these macros greatly simplifies signal handling for the application, and it is recommended that applications use these macros instead of directly invoking the methods of the `SignalManger` class, except for changing the set of handled signals.

If a signal does occur, process control jumps to the end of the signal block, and the `sigHandled()` method of the signal manager can be called. The application may need to have the same statements inside the `sigHandled()` check as those outside of the signal handling block. For example, if a file needs to be closed before the end of the block, the same call to the close function must be made within the `sigHandled()` check. Careful application design can reduce the amount of code replication, however.

Listing 5.2 shows how an application can indicate what signals to handle. In this example, only the `SIGUSR1` signal would be handled.

Listing 5.2: Specifying Signals to the **SignalManger**

```
1  #include <be_error_signal_manager.h>
2  using namespace BiometricEvaluation;
3
4  int main(int argc, char *argv[])
5  {
6      Error::SignalManager *sigmgr = new Error::SignalManager();
7
8      sigset_t sigset;
9      sigemptyset(&sigset);
10     sigaddset(&sigset, SIGUSR1);
11     sigmgr->setSignalSet(sigset);
12
13     FILE *fp = fopen( ... );
14     BEGIN_SIGNAL_BLOCK(sigmgr, sigblock2);
15         // code that may result in signal generation
16         fclose(fp);
17     END_SIGNAL_BLOCK(asigmgr, sigblock2);
```

```
18      if (sigmgr->sigHandled()) {
19          cout << "SIGUSR1 occurred." << endl;
20          fclose(fp);
21      }
22  }
```

Chapter 6

Input/Output

The `IO` package is used by applications for the common types of input and output: managing stores of data, log files, and individual file management. The goal of using the `IO` API is to relieve applications of the need to manage low-level I/O operations such as file opening, writing, and error handling. Furthermore, by using the classes defined in `IO`, the actual storage mechanism used for data can be managed efficiently and placed in a consistent location for all applications.

Many classes manage persistent storage within the file system, taking care of file open and close operations, as well as error handling. When errors do occur, exceptions are thrown, which then must be handled by the application.

6.1 Utility

The `IO::Utility` namespace provides functions that are used to manipulate the file system and other low-level mechanisms. These functions can be used by applications in addition to being used by other classes within the Biometric Evaluation framework. The functions in this package are used to directly manipulate objects in the POSIX file system, or to check whether a file object exists.

6.2 Record Management

The `IO::RecordStore` class provides an abstraction for performing record-oriented input and output to an underlying storage system. Each implementation of the `RecordStore` provides a self-contained entity to manage data on behalf of the application in a reliable, efficient manner.

Many biometric evaluations generate thousands of files in the form of processed images and biometric templates, in addition to consuming large numbers of files as input. In many file systems, managing large numbers of files in not efficient, and leads to longer run times as well as difficulty in backing up and processing these files outside of the actual evaluation.

The `RecordStore` abstraction de-couples the application from the underlying storage, enabling the implementation of different strategies for data management. One simple strategy is to store each record into a separate file, reproducing what has typically been done in the evaluation software itself. Archive files and small databases are other implementation strategies that have been used.

Use of the `RecordStore` abstraction allows applications to switch storage strategy by changing a few lines of code. Furthermore, error handling is consistent for all strategies by the use of common exceptions.

`RecordStore`s provide no semantic meaning to the nature of the data that passes through the store. Each record is an opaque object, given to the store as a pointer and data length, and is associated with a string the which is the key. Keys must be unique and are associated with a single record. Attempts to insert multiple records with the same key result in an exception being thrown.

Listing 6.1 illustrates the use of a database `RecordStore` within an application.

> **Listing 6.1: Using a `RecordStore`**

```
1  #include <iostream>
2  #include <be_io_dbrecstore.h>
3  int
4  main(int argc, char* argv[]) {
5
6      IO::DBRecordStore *rs;
7      try {
8          rs = new IO::DBRecordStore("myRecords", "My Record Store", "");
9      } catch (Error::Exception& e) {
10         cout << "Caught " << e.getInfo() << endl;
11         return (EXIT_FAILURE);
12     }
13     auto_ptr<IO::DBRecordStore> ars(rs);
14
15     try {
16         uint8_t *theData;
17
18         theData = getSomeData();
19         ars->insert("key1", theData);
20
21         theData = getSomeData();
22         ars->insert("key2", theData);
23
24     } catch (Error::Exception& e) {
25         cout << "Caught " << e.getInfo() << endl;
26         return (EXIT_FAILURE);
27     }
28
29     // Some more processing where new data for a key comes in ...
30     theData = getSomeData();
31     ars->replace("key1", theData);
32
33     // Obtain the data for all keys ...
34     string theKey;
35     while (true) {
36         uint64_t len = rs->sequence(theKey, theData);
37         cout << "Read data for key " << theKey << " of length " << len << endl;
38     }
39     // The data for the key is no longer needed ...
40     ars->remove("key1");
41 }
```

6.3 Logging

Many applications are required to log information during their processing. In particular, the evaluation test drivers often create a log record for each call to the software under test. There is a need for the log entries to be consistent, yet any logging facility must be flexible in accepting the type of data that is to be written to the log file.

The logging classes in the `IO` package provide a straight-forward method for applications to record their progress without the need to manage the low-level output details. There are two classes, `IO::LogCabinet`

and `IO::LogSheet` that are used to perform consistent logging of information by applications. A `LogCabinet` contains a set of `LogSheet` s.

A `LogSheet` is an output stream (subclass of `std::ostringstream`), and therefore can handle built-in types and any class that supports streaming. The example code in Listing 6.2 shows how an application can use a `LogSheet`, contained within a `LogCabinet`, to record operational information.

Log sheets are simple text files, with each entry numbered by the `LogSheet` class when written to the file. The description of the sheet is placed at the top of the file during construction of the *LogSheet* object. A call to the `newEntry()` method commits the current entry to the log file, and resets the write position to the beginning of the entry buffer.

In addition to streaming by using the `LogSheet::<<` operator, applications can directly commit an entry to the log file by calling the `write()` method, thereby not disrupting the entry that is being formed. After an entry is committed, the entry number is automatically incremented.

The example in Listing 6.2 shows application use of the logging facility.

> **Listing 6.2: Using a `LogSheet` within a `LogCabinet`**

```
1  #include <be_io_logcabinet.h>
2  using namespace BiometricEvaluation;
3  using namespace BiometricEvaluation::IO;
4
5  LogCabinet *lc;
6  try {
7      lc = new LogCabinet(lcname, "A Log Cabinet", "");
8  } catch (Error::ObjectExists &e) {
9      cout << "The Log Cabinet already exists." << endl;
10     return (-1);
11 } catch (Error::StrategyError& e) {
12     cout << "Caught " << e.getInfo() << endl;
13     return (-1);
14 }
15 auto_ptr<LogCabinet> alc(lc);
16 try {
17     ls = alc->newLogSheet(lsname, "Log Sheet in Cabinet");
18 } catch (Error::ObjectExists &e) {
19     cout << "The Log Sheet already exists." << endl;
20     return (-1);
21 } catch (Error::StrategyError& e) {
22     cout << "Caught " << e.getInfo() << endl;
23     return (-1);
24 }
25 ls->setAutoSync(true);   // Force write of every entry when finished
26 int i = ...
27 *ls << "Adding an integer value " << i << " to the log." << endl;
28 ls->newEntry();          // Forces the write of the current entry
29 .........
30 delete ls;
31 return;                  // The LogCabinet is destructed by the auto_ptr
```

6.4 Properties

The `Properties` class is used to store simple key-value string pairs, with the option to save to a file. Applications can use a `Properties` object to manage runtime settings that are persistent across invocations, or to simply store some settings in memory only.

Listing 6.3: Using a `Properties` Object

```
1  IO::Properties *props;
2  string fname = "test.prop";
3  try {
4      props = new IO::Properties(fname);
5  } catch (Error::StrategyError &e) {
6      cerr << "Caught " << e.getInfo()  << endl;
7      return;
8  } catch (Error::FileError& e) {
9      cerr << "A file error occurred: " << e.getInfo() << endl;
10     return;
11 }
12 props->setProperty("foo", "bar");
13 props->setProperty("theAnswer", "42");
14     :
15     :
16     :
17 try {
18     int64_t theAnswer = props->getProperty("theAnswer");
19     cout << "The answer is " << theAnswer << endl;
20 } catch (Error::ObjectDoesNotExist &e) {
21     cerr << "The answer is elusive." << endl;
22     return;
23 }
24 string fooProp = props->getProperty("foo");
25 cout << "Foo is set to " << fooProp << endl;
26     :
27     :
28     :
29 try {
30     props->removeProperty("foo");
31 } catch (Error::ObjectDoesNotExist &e) {
32     cerr << "Failed to remove property." << endl;
33 }
```

6.5   Compressor

Support for data compression and decompression can be found in the Biometric Evaluation Framework through the `Compressor` class hierarchy. `Compressor` is an abstract base class defining several pure-virtual methods for compression and decompression of buffers and files. Derived classes implement these methods and can be instantiated through the factory method in the base class. As such, children should also be enumaterated within `Compressor::Kind`. The Biometric Evaluation Framework comes with an example, `GZIP`, which compresses and decompresses the `gzip` format through interaction with `zlib` [4].

Listing 6.4: Using a `Compressor` Object

```
1  tr1::shared_ptr<IO::Compressor> compressor;
2  Memory::uint8Array compressedBuffer, largeBuffer = /* ... */;
3  try {
4          compressor = IO::Compressor::createCompressor(Compressor::Kind::GZIP);
5          /* Overloaded for all combination of buffer and file */
6          compressor->compress("largeInputFile", "compressedOutputFile");
7          compressedBuffer = compressor->compress(largeBuffer);
8  } catch (Error::Exception &e) {
```

```
 9           cerr << "Could not compress (" << e.getInfo() << ')' << endl;
10 }
```

Different `Compressor`s may be able to respond to options that tune their operations. These options (and approved values) should be well-documented in the child class, however, a no-argument constructor of a child `Compressor` should automatically set any required options to default values. Setting and retrieving these options is very similar to interacting with a `Properties` object (see Section 6.4 on page 17).

Listing 6.5: Setting `Compressor` Options

```
1 tr1::shared_ptr<IO::Compressor> compressor =
2     IO::Compressor::createCompressor(Compressor::Kind::GZIP);
3
4 /* A large GZIP chunk size can speed operations on systems with copious RAM */
5 compressor->setOption(IO::GZIP::CHUNK_SIZE, 32768);
```

Chapter 7

Time and Timing

The `Time` package within the Biometric Evaluation Framework provides a set of classes for performing timing-related operations, such as elapsed time and limiting execution time.

7.1 Elapsed Time

The `Timer` class provides applications a method to determine how long a block of code takes to execute. On many systems (e.g. Linux) the timer resolution is in microseconds.

Listing 7.1 shows how an application can use a `Timer` object to limit obtain the amount of time used for the execution of a block of code.

Listing 7.1: Using the `Timer`

```
1  #include <be_time_timer.h>
2
3  int main(int argc, char *argv[])
4  {
5          Time::Timer timer = new Time::Timer();
6
7          try {
8                  atimer->start();
9                  // do something useful, or not
10                 atimer->stop();
11                 cout << "Elapsed time: " << atimer->elapsed() << endl;
12         } catch (Error::StrategyError &e) {
13                 cout << "Failed to create timer." << endl;
14         }
15 }
```

7.2 Limiting Execution Time

The `Watchdog` class allows applications to control the amount of time that a block of code has to execute. The time can be *real* (i.e. "wall") time, or *process* time (not available on Windows). One typical usage for a `Watchdog` timer is when a call is made to a function that may never return, due to problems processing an input biometric image.

`Watchdog` timers can be used in conjunction with `SignalManager` in order to both limit the processing time of a call, and handle all signals generated as a result of that call. See 5.2 for information on the `SignalManager` class.

One restriction on the use of `Watchdog` is that the POSIX calls for signal management (`signal(3)`, `sigaction(2)`, etc.) cannot be invoked inside of the `WATCHDOG` block. This restriction includes calls to `sleep(3)` because it is based on signal handling as well.

Listing 7.2 shows how an application can use a `Watchdog` object to limit the about of process time for a block of code.

Listing 7.2: Using the **Watchdog**

```
1  #include <be_time_watchdog.h>
2  int main(int argc, char *argv[])
3
4      Time::Watchdog theDog = new Time::Watchdog(Time::Watchdog::PROCESSTIME);
5      theDog->setInterval(300);    // 300 microseconds
6
7      Time::Timer timer;
8
9      BEGIN_WATCHDOG_BLOCK(theDog, watchdogblock1);
10         timer.start();
11         // Do something that may take more than 300 usecs
12         timer.stop();
13         cout << "Total time was " <<  timer.elapsed() << endl;
14     END_WATCHDOG_BLOCK(theDog, watchdogblock1);
15     if (theDog->expired()) {
16         timer.stop();
17         cerr << "That took too long." << endl;
18     }
19 {
20 }
```

Within the `Watchdog` header file, two macros are defined: `BEGIN_WATCHDOG_BLOCK()` and `END_WATCHDOG_BLOCK()`, each taking the `Watchdog` object and label as parameters. The label must be unique for each `WATCHDOG` block. The use of these macros greatly simplifies `Watchdog` timers for the application, and it is recommended that applications use these macros instead of directly invoking the methods of the `Watchdog` class, except for setting the timeout value.

Any processing that is normally done at the end of the `WATCHDOG` block must also be done within the `expired()` check due to the fact that process control jumps to the end of the `WATCHDOG` block in the event of a timeout. A typical example is the use of the `Timer` object inside a `WATCHDOG` block, as the example in Listing 7.2 shows. In most cases, however, careful application design can remove the need for duplicate code. In the example, placing the `Timer` `start()`/`stop()` calls outside of the `WATCHDOG` block simplifies the coding, although the small amount of time for the `WATCHDOG` setup and tear down would be included in the time.

Chapter 8

Process Information

The `Process` package is a set of APIs used to gather information on a process, limit the capabilities of a process, and create manage processes.

8.1 Process Statistics

When a application is running, there is a need to obtain information of the process executing that application. The Process API can be used by the application itself to gather statistics related to the current amount of memory being used, the number of threads, and other items. Biometric evaluation test drivers are linked against a third party library, and therefore, the application writer does not control the thread count or memory usage for much of the processing. Listing 8.1 shows how an application can use the `Statistics` API.

Listing 8.1: Gathering Process Statistics

```cpp
#include <be_error_exception.h>
#include <be_process_statistics.h>
using namespace BiometricEvaluation;

int main(int argc, char *argv[])
{
    Process::Statistics stats;
    uint64_t userstart, userend;
    uint64_t systemstart, systemend;
    uint64_t diff;
    try {
        stats.getCPUTimes(&userstart, &systemstart);

        // Do some long processing....

        stats.getCPUTimes(&userend, &systemend);
        diff = userend - userstart;
        cout << "User time elapsed is " << diff << endl;
        diff = systemend - systemstart;
        cout << "System time elapsed is " << diff << endl;
    } catch (Error::Exception) {
        cout << "Caught " << e.getInfo() << endl;
    }

}
```

    In addition to using the `Process` API to gather statistics to be returned from the function call, the API provides a means to have a "standard" set of statistics logged either synchronously or asynchronously to a `LogSheet` (See Section 6.3 on page 16) contained within a `LogCabinet`. Applications can start and stop logging at will to this `LogSheet`. Post-mortem analysis can then be done on the entries in the `LogSheet`. Listing 8.2 shows the use of logging.

    The `LogSheet` will have a file name constructed from the process name (i.e. the application executable) and the process ID. An example `LogSheet` contains this information at the start:

```
Description: Statistics for test_be_process_statistics (PID 28370)
# Entry Usertime Systime RSS VMSize VMPeak VMData VMStack Threads
E0000000001 728889 6998 1788 57472 62612 31020 84 1
E0000000002 1300802 6998 1792 57472 62612 31020 84 1
```

    The `Statistics` object creates the `LogSheet` with an appropriate description and comment entry with column headers. Each gathering of the statistics results in a single log entry.

Listing 8.2: Logging Process Statistics

```cpp
#include <be_error_exception.h>
#include <be_io_logcabinet.h>
#include <be_process_statistics.h>
using namespace BiometricEvaluation;

int main(int argc, char *argv[])
{
    IO::LogCabinet lc("statLogCabinet", "Cabinet for Statistics", "");

    Process::Statistics *logstats;
    try {
        logstats = new Process::Statistics(&lc);
    } catch (Error::Exception &e) {
        cout << "Caught " << e.getInfo() << endl;
        return (EXIT_FAILURE);
    }
    try {
        while (some_processing_to_do) {
            // Do the work
            // Synchronously log after the work is done.
            logstats->logStats();
        }
    } catch (Error::Exception &e) {
        cout << "Caught " << e.getInfo() << endl;
        delete logstats;
        return (EXIT_FAILURE);
    }

    // Set up asynchronous logging, every second
    try {
        logstats->startAutoLogging(1);
    } catch (Error::ObjectExists &e) {
        cout << "Caught " << e.getInfo() << endl;
        delete logstats;
        return (EXIT_FAILURE);
    }

    // Do some other work
```

```
39
40      // Stop logging
41      logstats->stopAutoLogging();
42      delete logstats;
43  }
```

8.2 Process Management

During a biometric evaluation or other long-running CPU-bound task, it's beneficial to make efficient use of
all the hardware available on the system. If your application is running on a multi-core machine, why not
make use of more than one core? BECommon aims to simply this by abstracting the usage of `fork(2)` and
`libpthread` to run multiple instances of the same function simultaneously.

8.2.1 Manager

There are three class hierarchies involved in the abstraction. The `BiometricEvaluation::Process::Manager`
classes control the technique of process manipulation that will be used. BECommon provides two example
abstractions: `ForkManager` and `POSIXThreadManager`. When using `ForkManager`, new processes
will be created with `fork(2)`, with mediated access to these new processes through the `Manager`. Likewise,
`POSIXThreadManager` creates new POSIX threads. Because both of theses classes inherit from `Manager`,
it is as trivial as changing the `Manager` object type to change how the workload is parallelized.

8.2.2 Worker

In the application using a `Manager`, a `Worker` subclass must be implemented. An example `Worker` is shown
in Listing 8.3. The entry-point for a `Worker` is the `workerMain()` method, which must be implemented
by the client application. Although `workerMain()` takes no arguments, data may be transmitted into the
object through `WorkerController`'s (8.2.3) `setParameter()` method. Within the `Worker` instance,
the parameters are then retrieved with `getParameter()` when provided with the unique parameter name.

    A responsible `Worker` performs its operations as fast as it can, however, at any given time, the `Manager`
may ask the `Worker` to stop. It then becomes the *responsibility of the* `Worker` to stop as soon as possible.
The `Worker` is notified of the stop request through its `stopRequested()` method. Note that the `Manager`
does **not** force the `Worker` to stop, though prolonged work or cleanup in the `Worker` would likely produce
undesired results in the client application. As such, a responsible `Worker` checkpoints itself to prepare for
premature stops requested by the `Manager`. While it is important for `Worker` s to stop as soon as possible
after the request is received, it is also important not to leave work in an unsynchronized state. In Listing 8.3,
notice how the `Employee` must continue the interaction with the `Customer` before a stop request is handled,
even if the `Employee`'s shift has ended. Leaving the method before the `Customer`'s order has been delivered
would leave the `Customer` object in an unsafe state (hungry).

> **Listing 8.3: A Responsible `Worker` Implementation**

```
1  #include <cstdlib>
2  #include <tr1/memory>
3  #include <queue>
4
5  #include <restaurant.h>
6
7  #include <be_process_forkmanager.h>
8
9  using namespace std;
10 using namespace BiometricEvaluation;
```

```
11  using namespace Restaurant;
12
13  class ResponsibleEmployeeTask : public Process::Worker
14  {
15  public:
16          int32_t
17          workerMain()
18          {
19                  int32_t status = EXIT_FAILURE;
20
21                  /* Retrieve objects assigned to this Task */
22                  tr1::shared_ptr<Employee> employee =
23                      tr1::static_pointer_cast<Employee>(
24                      this->getParameter("employee"));
25                  tr1::shared_ptr< queue<Customer*> > customers =
26                      tr1::static_pointer_cast< queue<Customer*> >(
27                      this->getParameter("customers")
28
29                  employee->clockIn();
30
31                  Customer *customer;
32                  /* Checkpoint after each customer */
33                  while (this->stopRequested() == false ||
34                      employee->isShiftOver() == false) {
35                          customer = customers->front();
36
37                          if (customer != NULL) {
38                                  employee->takeOrder(customer);
39                                  employee->cookFood(customer);
40                                  employee->deliverOrder(customer);
41
42                                  customers->pop();
43                          }
44                  }
45
46                  employee->settleCashDrawer();
47                  employee->clockOut();
48
49                  status = EXIT_SUCCESS;
50                  return (status);
51          }
52          ~ResponsibleEmployeeTask() {}
53  };
```

After a `Manager` starts its `Worker`s, the `Manager` has the option of waiting until all `Worker`s exit `workerMain()` before continuing code execution. If not waiting, there are several methods the `Manager` can perform to keep track of the status of the `Worker`s. Even if not waiting for `Worker`s to return, a responsible `Manager` will wait a reasonable amount of time for `Worker`s to `return` before application termination. An example of this reasonable waiting period can be seen in Listing .

8.2.3 WorkerController

The final piece of the process management puzzle is the `WorkerController` hierarchy. This class decorates and mediates communication between the `Manager` and the `Worker`. `WorkerController` objects may only be instantiated by a `Manager` object. All communications to the `Worker` (e.g. `isWorking()`) should be delegated through the `WorkerController`. If defining a new `Manager`, note that the `Worker`

Controller may seem unnecessary for the parallelization technique being employed. It's true that some parallelization techniques may not require this "middle-man" approach, but others do. Do not be concerned if a WorkerController implementation ends up being nothing more than a "pass-thru" to the Worker.

Listing 8.4 is a continuation of Listing 8.3 on page 25 demonstraiting the use of Manager s and Worker Controller s.

Listing 8.4: Using Manager s and WorkerController s

```
1  int
2  main(
3      int argc,
4      char *argv[])
5  {
6          static const uint32_t numEmployees = 3;
7          int status = EXIT_FAILURE;
8
9          tr1::shared_ptr<Process::Manager> shiftLeader(new Process::ForkManager);
10         queue<Customer*> *customers = new queue<Customer*>();
11
12         /* Create Employees (Workers/WorkerControllers) */
13         tr1::shared_ptr<Process::WorkerController> employees[numEmployees];
14         for (uint32_t i = 0; i < numEmployees; i++) {
15                 employees[i] = shiftLeader->addWorker(
16                     tr1::shared_ptr<ResponsibleEmployeeTask>(
17                     new ResponsibleEmployeeTask()));
18
19                 /* Assign employees to each Task */
20                 employees[i]->setParameter("employee",
21                     tr1::shared_ptr<Employee>(new Employee()));
22                 employees[i]->setParameter("customers",
23                     tr1::shared_ptr< queue<Customer*> >(customers);
24         }
25
26         /* Employees start serving customers while shift leader manages */
27         shiftLeader->startWorkers(false);
28
29         /* Customers enter the queue... */
30         queue<Restaurant::AdministrativeTasks> adminTasks;
31         adminTasks.push("Inventory");
32         adminTasks.push("Customer Complaints");
33         adminTasks.push("Clean Dining Room");
34
35         while (shiftLeader->getNumActiveWorkers() != 0) {
36                 shiftLeader->doTask(adminTasks.front());
37                 adminTasks.pop();
38         }
39
40         /* ...end of the day */
41         for (uint32_t i = 0; i < numEmployees; i++)
42                 if (employees[i]->isWorking())
43                         shiftLeader->stopWorker(employees[i]);
44
45         /*
46          * Wait a reasonable amount of time before locking up for the night
47          * (in this case, indefinitely).
48          */
```

```
49          while (shiftLeader->getNumActiveWorkers() > 0)
50              sleep(1);
51
52          shiftLeader->armAlarmAndExit();
53
54          status = EXIT_SUCCESS;
55          return (status);
56  }
```

8.2.4 Communications

`Manager` s and `Worker` s might have good reason to communicate arbitrary messages directly. A communications mechanism is built-in to the Process Management model to facilitate such communications. The type and content of the message is completely up to the client implementation, since messages are sent as `AutoArray` s. A `Manager` does not directly send messages to a `Worker`. This service is provided by the `WorkerController` (via `sendMessageToWorker()`).

`Manager` s can keep an eye on incoming messages by calling the (optionally blocking) `waitForMessage()` method. This method will return a handle to the `Worker` that sent a message. Alternatively, the `Manager` can invoke `getNextMessage()` (again, blocking optional) to immediately receive the next message.

Listing 8.5 and Listing 8.6 are continuations of Listing 8.3 on page 25 and Listing 8.4 on the previous page respectively, showing an example of communication, using `std::string` messages.

Listing 8.5: `Worker` Communication

```
1       Memory::uint8Array msg;
2
3       /* Deal with next customer unless Manager interrupts in next second */
4       if (this->waitForMessage(1)) {
5               if (this->getMessageFromManager(msg)) {
6                       Action action = Restaurant::messageToAction(msg);
7                       switch (action) {
8                       case TAKE_BREAK:
9                               employee->goOnBreak();
10                              break;
11                      /* ... */
12                      }
13              }
14      }
15
16      /* ... */
17
18      if (customer->isComplaining()) {
19              sprintf((char *)&(*msg), "Customer Complant");
20              this->sendMessageToManager(msg);
21      }
```

Listing 8.6: `Manager` Communication

```
1       tr1::shared_ptr<Process::WorkerController> sender;
2       Memory::uint8Array msg;
3
4       /* Do routine tasks unless employee has concern in the next 2 seconds */
5       while (this->getNextMessage(sender, msg, 2)) {
```

```
 6              Action action = Restaurant::messageToAction(msg);
 7              switch (action) {
 8              case CUSTOMER_COMPLAINT:
 9                      sprintf((char *)&(*msg), "I'll take care of it.");
10                      this->sendMessageToWorker(msg);
11                      break;
12              /* ... */
13              }
14      }
15
16      /* ... */
17
18      /*  Closing Time */
19      sprintf((char *)&(*msg), "Clock out and go home.");
20      this->broadcastMessage(msg);
```

Chapter 9

System

The System package provides a set of functions in the that return information about the hardware and operating system. This information can be used by applications to determine the amount of real memory, number of central processing units, or current load average. This information can be used to dynamically tailor the application behavior, or simply to provide additional information in a runtime log.

Listing 9.1 shows how an application can spawn several child processes based on the number of CPUs and memory available. Note that this information may not be available on all platforms, and therefore, the application must be prepared to handle that situation.

Listing 9.1: Using the **System** CPU Count Information

```
1  #include <iostream>
2  #include <be_system.h>
3
4  using namespace BiometricEvaluation;
5
6  int
7  main(int argc, char* argv[]) {
8
9      // perform some application setup ...
10
11     uint32_t cpuCount;
12     uint64_t memSize, vmSize;
13     try {
14         cpuCount =  System::getCPUCount();
15         cpuCount--;    // subtract one CPU for the parent process
16         memSize = System::getRealMemorySize();
17         Process::Statistics::getMemorySizes(NULL, &vmSize, NULL, NULL, NULL);
18         memSize -= vmSize;   // subtract off memory used by parent
19
20         // Give each child a fraction of the memory
21         spawnChildren(cpuCount, memSize / cpuCount);
22     } catch (Error::NotImplemented) {
23             cout << "Running a single process only." << endl;
24     }
25
26     // processing done by parent ...
27 }
```

Chapter 10

Image

The `Image` package maintains the classes and other information related to images and image processing. Within the Biometric Evaluation Framework, many classes refer to images, such as when dealing with fingerprint data. Many biometric data standards supply the actual image encoded in one of several standard formats. Applications can retrieve the image as stored in the record, or decompressed by the Image class into a "raw" format. Therefore, within the BECommon, several of the common compression formats are supported, removing the need for applications to decompress the image directly, while maintaining access to the as-recorded image format.

10.1   The Image Namespace

The `Image` namespace contains several data types used to represent aspects of an image. The types defined are chiefly used to retrieve common information from images stored in an `Image` class (section 10.2). Data types in the Image namespace do not perform any translation of scale units or sizing, as each set of attributes is copied directly from the image data itself when possible.

The same applies to images encapsulated in biometric records. Although some biometic records have fields for image attributes like dimensions and resolution, the corresponding fields of an `Image` class are **not** populated with their contents. The `Image` namespace data types *are* used outside of the namespace, such as in finger views, to retrieve image attributes stored as part of the biometric record. Applications can compare those values against the values within the `Image` object, as in most cases those values are taken directly from the underlying image data. See Chapter 14 on page 45 for more information on image-based biometric records.

The `Image` namespace contains all of the `Image` classes that are used to represent an image. These classes are described in the following sections.

10.2   The Image Class

The `Image` class is an abstract base class that defines a set of minimum functionality for all supported image formats. Once an `Image` has been constructed, it may not be modified. For any supported image format, the following information is required to be accessible:

- Original binary data

- Compression algorithm

- Decompressed ("raw") format binary data (grayscale, full color)

- Depth

- Dimensions (width, height)

- Resolution (horizontal, vertical)

A rudimentary implementation of generating a grayscale image is provided by the `Image` class in `getRaw GrayscaleData()`. This implementation calculates the luminance value Y (of YCbCr) for each pixel of a color image. The resulting image always uses 8-bits to represent a pixel, but can return a raw image using 2 gray levels (1-bit) or 256 gray levels (8-bit). The 1-bit algorithm quantizes to black when the 8-bit color value is ≤127. `Image` subclasses may override and implement their own grayscale conversion methods.

Also of interest in the Image class is `valueInColorspace()`, a static function to convert color values between bit depths.

10.3 Raw Image

The `RawImage` class represents a decompressed image, or an image where `getRawData()` would return the exact same data as `getData()`. `RawImage` has no special implementation or additional methods.

10.4 JPEG

The `JPEG` class represents an image encoded according to the JPEG image standard [8]. Decompression and grayscale conversion are accomplished via `libjpeg` [6].

As of version 8.0, `libjpeg` provided a way to handle JPEG images existing within in-memory buffers, as opposed to on-disk files. Because the Image class requires in-memory buffers, `JPEG` includes a JPEG memory source manager implementation, but it is built only if a version of `libjpeg` older than 8.0 is detected at compile-time.

`JPEG` provides a static function to determine whether or not a data buffer appears to be encoded in the JPEG image standard format. Errors within libjpeg will be caught and rethrown as `Exception` s.

10.5 JPEGL

Similar to `JPEG`, the `JPEGL` class performs `Image` class services for lossless JPEG encoded images. JPEGL decompression is performed by NIST Biometric Image Software 's `libjpegl` [12].

10.6 JPEG2000

The `JPEG2000` class provides Image class functionality to JPEG 2000-encoded images [7]. The class makes an attempt to support the following JPEG 2000 codecs:

- JPEG 2000 codestream (.j2k)

- JPEG 2000 compressed image data (.jp2)

- JPEG 2000 interactive protocol (.jpt)

Decompression is provided by the OpenJPEG library (`libopenjpeg`) [11]. `JPEG2000` also provides a static function to test whether or not an image appears to be JPEG 2000-encoded.

Not all information required by the `Image` class is present in a JPEG 2000-encoded image. In particular, some codecs and encoders omit the "Display Resolution Box." It is generally accepted that the resolution will be 72 pixels-per-inch when the "Display Resolution Box" is not present.

Errors within `libopenjpeg` will be caught and rethrown as `Exception` s.

10.7   NetPBM

The `NetPBM` class provides Image class functionality to all types of NetPBM formatted images, up to 48-bit depth. This includes the following formats:

- ASCII Portable Bitmap (P1, .pbm)

- ASCII Portable Graymap (P2, .pgm)

- ASCII Portable Pixmap (P3, .ppm)

- Binary Portable Bitmap (P4, .pbm)

- Binary Portable Graymap (P5, .pgm)

- Binary Portable Pixmap (P6, .ppm)

`NetPBM` provides some of its more general use parsing algorithms as static functions for use outside of the class. This includes ASCII to binary pixel conversion. A function to test for NetPBM formats is also provided.

10.8   PNG

The `PNG` class represents an image encoded according to the PNG image standard [5]. Decompression is provided by `libpng` [14].

`PNG` provides a static function to test whether or not an image appears to be encoded in the PNG image standard format. Errors within `libpng` are caught and rethrown as `Exception`s.

10.9   WSQ

Images encoded in the WSQ-image standard [16] are represented by the `WSQ` class. The WSQ decompressor found in NIST Biometric Image Software [12], `libwsq`, is used by this class. The class provides a static function to determine whether or not an image appears to be encoded in the WSQ format.

Errors from the `libwsq` will be displayed through `stderr` and will **not** be rethrown as `Exception`s.

Chapter 11

Text

The `Text` package consists of functions to perform common operations on `string`s and `char` arrays. Many of the operations may be considered "trivial," but are used often enough within the Biometric Evaluation Framework and other applications that a common implementation in BECommon is more than warranted. A complete listing of functions is available in the documentation appendix for `BiometricEvaluation::Text`2.

Listing 11.1 shows how to use the `split()` function from the `Text` package. `split()` can separate a `string` into tokens delimited by a character, useful for processing comma- or space-separated text files (such files could be produced by a `LogSheet` (Section 6.3 on page 16), for instance). Here, a text file containing metadata for an image is being parsed, perhaps to be passed to the `RawImage` constructor (Section 10.3 on page 34).

Listing 11.1: Tokenizing a `string`

```
1  /* Definition of input strings */
2  static const vector<string>::size_type filenameToken = 0;
3  static const vector<string>::size_type widthToken = 1;
4  static const vector<string>::size_type heightToken = 2;
5  static const vector<string>::size_type depthToken = 3;
6
7  /* Split the string, presumably input from a file */
8  string input = "/mnt/raw\\ images/1.raw 500 500 8";
9  vector<string> tokens = Text::split(input, ' ', true);
10
11 /* Assign the retrieved tokens */
12 string filename;
13 uint32_t width, height, depth;
14 try {
15         filename = tokens.at(filenameToken);    /* "/mnt/raw images/1.raw" */
16         width = atoi(tokens.at(widthToken).c_str());    /* "500" */
17         height = atoi(tokens.at(heightToken).c_str());  /* "500" */
18         depth = atoi(tokens.at(depthToken).c_str());    /* "8" */
19 } catch (out_of_range) {
20         throw Error::FileError("Malformed input");
21 }
```

Notice the `true` parameter to `split()` in Listing 11.1. This instructs `split()` to not tokenize based on an escaped delimiter. If `false`, the first token would be split into two at the presence of the delimiter.

`Text` also contains functions to perform hashing via `OpenSSL`. A two-line program that emulates the command-line `md5sum` program is shown in Listing 11.2. Changing the digest parameter to `"sha1"` would make the program emulate `openssl sha1`.

Listing 11.2: md5sum via BECommon

```
1  #include <cstdlib>
2  #include <iostream>
3
4  #include <be_io_utility.h>
5  #include <be_text.h>
6  #include <be_memory_autoarray.h>
7
8  using namespace std;
9  using namespace BiometricEvaluation;
10
11 int
12 main(
13     int argc,
14     char *argv[])
15 {
16         if (argc == 0)
17                 return (EXIT_FAILURE);
18
19         try {
20                 Memory::uint8Array file = IO::Utility::readFile(argv[1]);
21                 cout << Text::digest(file, file.size(), "md5") << "  " <<
22                     argv[1] << endl;
23         } catch (Error::Exception) {
24                 return (EXIT_FAILURE);
25         }
26
27         return (EXIT_SUCCESS);
28 }
```

Chapter 12

Feature

The `Feature` package contains those items that relate to the representation of biometric features, such as fingerprint minutiae, facial features (eyes, etc.), and related information. Objects of these class types are typically associated with `View` (Chapter) or `DataInterchange` (Chapter) objects. For example, a minutiae object is usually obtained from a finger view, which may have been obtained from a data interchange object representing an entire biometric record for an individual.

The data contained within a `Feature` object is represented as the "native" format as it was extracted from the underlying data record. There is no translation to a common format and it is the application's responsibility to interpret or translate the data as necessary.

Currently, fingerprint and palm print minutiae are the features supported within the BECommon. As development continues, additional features contained within biometric data records will be supported.

12.1 ANSI/NIST Features

The ANSI/NIST [3] standard defines several features represented as data elements within a record. Fingerprint and palm minutiae is contained within Type-9 record. The `AN2K7Minutiae` class, contained in the Feature package, represents a single Type-9 record. An object of this class can be constructed directly from a complete ANSI/NIST record. However, it is more common for an application to retrieve these objects from the `AN2KView` object defined in the `Finger` package (Chapter).

See Listing for a complete example of how to obtain the fingerprint minutiae data from an ANSI/NIST record.

12.2 ISO/INCITS Features

The ISO [2] and INCITS [1] fingerprint minutiae standards are represented within BECommon with the same class, `INCITSMinutiae`, as the minutiae format is identical in both standards.

Listing shows how to create a view object for the fingerprint minutiae record contained in a file.

Chapter 13

Finger

One of the most commonly used biometric source is the fingerprint. Multiple types of information can be derived from a fingerprint, including minutiae and the pattern, such as whorl, etc. The `Finger` package contains the types, classes, and other items that are related to fingers and fingerprints. Objects of the `Finger` classes are typically not used in a stand-alone fashion, but are usually obtained from an object in the `DataInterchage` (Chapter 15 on page 47) package.

Several enumerated types are defined in the `Finger` package. The types are used to represent those elements related to fingers and fingerprints that are common across all data formats. Types that represent finger position, impression type, and others are included in the package. Stream operators are defined for these types so they can be printed in human-readable format.

Most of the classes in the `Finger` package represent data taken directly from a record in a standard format (e.g. ANSI/NIST [3]). In addition to general information, such as finger position, other information may be represented: The source of the finger image; the quality of the image, etc. In addition to this descriptive information, the finger object will provide the set of derived minutiae or other data sets.

When representing the information about a finger (and fingerprint), the class in the `Finger` package implements the interface defined in the `View` package. A finger is a specific type of view in that it represents all the available information about the finger, including the source image, minutiae (often in several formats), as well as the capture data (date, location, etc.)

13.1 ANSI/NIST Minutiae Data Record

Finger views are objects that represent all the available information for a specific finger as contained in one or more biometric records. For example, an ANSI/NIST file may contain a Type-3 record (finger image) and an associated Type-9 record (finger minutiae). A finger view object based on the ANSI/NIST record can be instantiated and used by an application to retrieve all the desired information, including the source finger image. The internals of record processing and error handling are encapsulated within the class.

The BECommon provides several classes that are derived from a base View class, contained within the `Finger` package. See Chapter 13 for more information on the types associated with fingers and fingerprints. This section discusses finger views, the classes which are derived from the general `View` class. These subclasses represent specific biometric file types, such as ANSI/NIST or INCITS/M1. In the latter case, two files must be provided when constructing the object because INCITS finger image and finger minutiae records are defined in two separate standards.

13.1.1 ANSI/NIST Finger Views

An ANSI/NIST record may contain one or more finger views, each based on a type of finger image. These Type-3, Type-4, etc. records contain the image and Type-9 minutiae data, among other information. These

record types are grouped into either the fixed- or variable-resolution categories, and are represented as specific classes within BECommon, AN2KViewFixedResolution and AN2KViewVariableResolution.

The AN2KMinutiaeDataRecord class represents all of the information taken from a ANSI/NIST Type-9 record. A Type-9 record may include minutiae data items in several formats (standard and proprietary) and the impression type code.

Listing 13.1 shows how an application can use the AN2KViewFixedResolution to retrieve image information, image data, and derived minutiae information from a file containing an ANSI/NIST record with Type-3 (fixed resolution image) and Type-9 (fingerprint minutiae) records.

Listing 13.1: Using an AN2K Finger View

```
1  #include <fstream>
2  #include <iostream>
3  #include <be_finger_an2kview_fixedres.h>
4  using namespace std;
5  using namespace BiometricEvaluation;
6
7  int
8  main(int argc, char* argv[]) {
9
10     Finger::AN2KViewFixedResolution *_an2kv
11     try {
12         _an2kv = new Finger::AN2KViewFixedResolution("type9-3.an2k",
13             TYPE_3_ID, 1);
14     } catch (Error::DataError &e) {
15         cerr << "Caught " << e.getInfo()  << endl;
16         return (EXIT_FAILURE);
17     } catch (Error::FileError& e) {
18         cerr << "A file error occurred: " << e.getInfo() << endl;
19         return (EXIT_FAILURE);
20     }
21     std::auto_ptr<Finger::AN2KView> an2kv(_an2kv);
22
23     cout << "Image resolution is " << an2kv->getImageResolution() << endl;
24     cout << "Image size is " << an2kv->getImageSize() << endl;
25     cout << "Image depth is " << an2kv->getImageDepth() << endl;
26     cout << "Compression is " << an2kv->getCompressionAlgorithm() << endl;
27     cout << "Scan resolution is " << an2kv->getScanResolution() << endl;
28
29     // Save the finger image to a file.
30     tr1::shared_ptr<Image::Image> img = an2kv->getImage();
31     if (img.get() == NULL) {
32         cerr << "Image was not present." << endl;
33         return (EXIT_FAILURE);
34     }
35     string filename = "rawimg";
36     ofstream img_out(filename.c_str(), ofstream::binary);
37     img_out.write((char *)&(img->getRawData()[0]),
38         img->getRawData().size());
39     if (img_out.good())
40             cout << "\tFile: " << filename << endl;
41     else {
42         img_out.close();
43         cerr << "Error occurred when writing " << filename << endl;
44         return (EXIT_FAILURE);
45     }
```

```
46      img_out.close();
47
48      // Get the finger minutiae sets. AN2K records can have more than one
49      // set of minutiae for a finger.
50
51      vector<Finger::AN2KMinutiaeDataRecord> mindata = an2kv->getMinutiaeDataRecordSet();
52  }
```

13.1.2   ISO/INCITS Finger Views

The ISO [10] and INCITS [9] standards typically use separate files for the source biometric data and the derived data. For example, the ISO 19794-2 standard is for fingerprint minutiae data, while 19794-4 is for finger image data. The corresponding BECommon view objects are constructed with both files, although a view can be constructed with only one file. In the latter case, the view object will represent only that information contained in the single file.

Listing 13.1 on the preceding page shows how an application can create a view from a ANSI/INCTIS 378 finger minutiae format record [1].

> **Listing 13.2: Using an INCITS Finger View**

```
 1  #include <stdlib.h>
 2  #include <fstream>
 3  #include <iostream>
 4  #include <be_finger_ansi2004view.h>
 5  #include <be_feature_incitsminutiae.h>
 6  using namespace std;
 7  using namespace BiometricEvaluation;
 8
 9  int
10  main(int argc, char* argv[]) {
11
12      Finger::ANSI2004View fngv;
13      try {
14          fngv = Finger::ANSI2004View("test_data/fmr.ansi2004", "", 3);
15      } catch (Error::DataError &e) {
16          cerr << "Caught " << e.getInfo()  << endl;
17          return (EXIT_FAILURE);
18      } catch (Error::FileError& e) {
19          cerr << "A file error occurred: " << e.getInfo() << endl;
20          return (EXIT_FAILURE);
21      }
22      cout << "Image resolution is " << fngv.getImageResolution() << endl;
23      cout << "Image size is " << fngv.getImageSize() << endl;
24      cout << "Image depth is " << fngv.getImageDepth() << endl;
25      cout << "Compression is " << fngv.getCompressionAlgorithm() << endl;
26      cout << "Scan resolution is " << fngv.getScanResolution() << endl;
27
28      Feature::INCITSMinutiae fmd = fngv.getMinutiaeData();
29      cout << "Minutiae format is " << fmd.getFormat() << endl;
30      Feature::MinutiaPointSet mps = fmd.getMinutiaPoints();
31      cout << "There are " << mps.size() << " minutiae points:" << endl;
32      for (size_t i = 0; i < mps.size(); i++)
33          cout << mps[i];
34
35          Feature::RidgeCountItemSet rcs = fmd.getRidgeCountItems();
```

43

```
36      cout << "There are " << rcs.size() << " ridge count items:" << endl;
37      for (int i = 0; i < rcs.size(); i++)
38          cout << "\t" << rcs[i];
39
40      Feature::CorePointSet cores = fmd.getCores();
41      cout << "There are " << cores.size() << " cores:" << endl;
42      for (int i = 0; i < cores.size(); i++)
43          cout << "\t" << cores[i];
44
45      Feature::DeltaPointSet deltas = fmd.getDeltas();
46      cout << "There are " << deltas.size() << " deltas:" << endl;
47      for (int i = 0; i < deltas.size(); i++)
48          cout << "\t" << deltas[i];
49
50      exit (EXIT_SUCCESS);
51  }
```

Chapter 14

View

Within the Biometric Evaluation Framework a `View` represents all the information that was derived from an image of a biometric sample. For example, with a fingerprint image, any minutiae that were extracted from that image, as well as the image itself, are contained within a single `View` object. In many cases the image may not be present, however the image size and other information is contained within a biometric record, along with the derived information. A `View` is used to represent these records as well.

In the case where a raw image is part of the biometric record, the `View` object's related `Image` (Chapter 10 on page 33 object will have identical size, resolution, etc. values because the `View` class sets the `Image` attributes directly. For other image types (e.g. `JPEG`) the `Image` object will return attribute values taken from the image data.

`View` s are high-level abstractions of the biometric sample, and concrete implementations of a `View` include finger, face, iris, etc. views based on a specific type of biometric. Therefore, `View` objects are not created directly, Subclasses, such as finger views (see Chapter 13 on page 41), represent the specific type of biometric sample.

Objects are created with information taken from a biometric data record, an ANSI/NIST 2007 file, for example. Most record formats contain information about the image itself, such as the resolution and size. The `View` object can be used to retrieve this information. However, the data may differ from that contained in the image itself, and applications can compare the corresponding values taken from the `Image` object (when available) to those taken from the `View` object.

Chapter 15

Data Interchange

The `DataInterchange` package consists of classes and other elements used to process an entire biometric data record, or set of records. For example, a single ANSI/NIST record, consisting of many smaller records (fingerprint images, latent data, etc.) can be accessed by instantiating a single object. Classes in this package typically use has-a relationships to classes in the `Finger` and other packages that process individual biometric samples.

The design of classes in the `DataInterchange` package allows applications to create a single object from a biometric record, such as an ANSI/NIST file. After creating this object, the application can retrieve the needed information (such as `Finger View` s Chapter 13 on page 41) from this object. A typical example would be to retrieve all images from the record and pass them into a function that extracts a biometric template or some other image processing.

15.1   ANSI/NIST Data Records

The ANSI/NIST Data Interchange package contains the classes used to represent ANSI/NIST [3] records. One class, `AN2KRecord`, is used to represent the entire ANSI/NIST record. An object of this class will contain objects of the `Finger` classes, as well as other packages. By instantiating the `AN2KRecord` object, the application can retrieve all the information and images contained in the ANSI/NIST record.

The `AN2KMinutiaeDataRecord` class represents an entire Type-9 record from an ANSI/NIST file. However, some components of this class are represented by classes in other packages. For example, the `AN2K7Minutiae` class in the `Feature` package represents the "standard" format minutiae in the Type-9 record.

Listing 15.1 shows how an application can retrieve all finger captures (Type-4 records) from an ANSI/NIST record. Once the Views are retrieved, the application obtains the set of minutiae records associated with that View.

Listing 15.1: Retrieving ANSI/NIST Finger Captures

```
1  #include <iostream>
2  #include <be_error_exception.h>
3  #include <be_finger_an2kview_capture.h>
4
5  int
6  main(int argc, char* argv[])
7  {
8      /*
9       * Call the constructor that will open an existing AN2K file and
10      * retrieve the first finger capture (Type-14) record.
11      */
```

```
12      std::auto_ptr<Finger::AN2KViewCapture> an2kv;
13      try {
14          an2kv.reset(new Finger::AN2KViewCapture("type9-14.an2k", 1));
15      } catch (Error::DataError &e) {
16          cout << "Caught " << e.getInfo()  << endl;
17          return (EXIT_FAILURE);
18      } catch (Error::FileError& e) {
19          cout << "A file error occurred: " << e.getInfo() << endl;
20          return (EXIT_FAILURE);
21      }
22
23      cout << "Get the set of minutiae data records: ";
24      vector<Finger::AN2KMinutiaeDataRecord> records =
25          an2kv->getMinutiaeDataRecordSet();
26      cout << "There are " << records.size() << " minutiae records." << endl;
27
28      /*
29       * Get the info from the first minutiae record in the View.
30       */
31      DataInterchange::AN2KMinutiaeDataRecord type9 = records[0];
32
33      /*
34       * Get the "standard" set of minutiae.
35       */
36      Feature::AN2K7Minutiae an2k7m = type9.getAN2K7Minutiae();
37
38      /*
39       * Obtain the minutiae points, ridge counts, cores, and deltas.
40       */
41      Feature::MinutiaPointSet mps;
42      Feature::RidgeCountItemSet rcs;
43      Feature::CorePointSet cps;
44      Feature::DeltaPointSet dps;
45      try {
46          mps = an2k7m->getMinutiaPoints();
47          rcs = an2k7m->getRidgeCountItems();
48          cps = an2k7m->getCores();
49          dps = an2k7m->getDeltas();
50
51      } catch (Error::DataError &e) {
52          cout << "Caught " << e.getInfo()  << endl;
53          return (EXIT_FAILURE);
54      }
55
56      cout << "There are " << mps.size() << " minutiae points:" << endl;
57      /*
58       * Print out the minutiae points.
59       */
60      for (int i = 0; i < mps.size(); i++) {
61          printf("(%u,%u,%u)\n", mps[i].coordinate.x, mps[i].coordinate.y,
62              mps[i].theta);
63      }
64      cout << "There are " << rcs.size() << " ridge counts:" << endl;
65      for (int i = 0; i < rcs.size(); i++) {
66          printf("(%u,%u,%u)\n", rcs[i].index_one, rcs[i].index_two,
67          rcs[i].count);
```

```
68        }
69        cout << "There are " << cps.size() << " cores." << endl;
70        cout << "There are " << dps.size() << " deltas." << endl;
71
72        cout << "Fingerprint Reader: " << endl;
73        try { cout << an2k7m->getOriginatingFingerprintReadingSystem() << endl; }
74        catch (Error::ObjectDoesNotExist) { cout << "<Omitted>" << endl; }
75
76        cout << "Pattern (primary): " <<
77        Feature::AN2K7Minutiae::convertPatternClassification(
78        an2k7m->getPatternClassificationSet().at(0)) << endl;
79
80        return(EXIT_SUCCESS);
81 }
```

Listing 15.2 shows how an application can retrieve all latent finger images from a set of ANSI/NIST record retrieved from a `RecordStore`. Using the `Image` object, the image's "raw" data can be retrieved and passed to another function for processing. Note that the image data may be stored in a compressed format inside the ANSI/NIST record, but is converted to raw format by the `Image` object.

Listing 15.2: Retrieving ANSI/NIST Latent Records

```
1  #include <be_io_recordstore.h>
2  #include <be_data_interchange_an2k.h>
3  using namespace BiometricEvaluation;
4
5  void
6  processImageData(uint8_t *buf, uint32_t size)
7  {
8        :
9        :
10       :
11       :
12 }
13
14 int
15 main(int argc, char* argv[]) {
16
17       std::tr1::shared_ptr<IO::RecordStore> rs;
18       try {
19           rs = IO::RecordStore::openRecordStore(rsname, datadir, IO::READONLY);
20       } catch (Error::Exception &e) {
21           cerr << "Could not open record store: " << e.getInfo() << endl;
22           return (EXIT_FAILURE);
23       }
24
25       /*
26        * Read some AN2K records and construct the View objects.
27        */
28       Utility::uint8Array data;
29       string key;
30       while (true) {             // Loop through all records in store
31           uint64_t rlen;
32           try {
33               rlen = rs->sequence(key, NULL);
34           } catch (Error::ObjectDoesNotExist &e) {
35               break;
```

```
36          } catch (Error::Exception &e) {
37              cout << "Failed sequence: " << e.getInfo() << endl;
38              return (EXIT_FAILURE);
39          }
40          data.resize(rlen);
41          try {
42              rs->read(key, data);
43              DataInterchange::AN2KRecord an2k(data);
44              std::vector<Finger::AN2KViewLatent> latents = an2k.getFingerLatents();
45              for (int i = 0; i < latents.size(); i++) {
46                  tr1::shared_ptr<Image::Image> img = latents[i].getImage();
47                  if (img != NULL) {
48                      cout << "\tCompression: " << img->getCompressionAlgorithm() << endl;
49                      cout << "\tDimensions: " << img->getDimensions() << endl;
50                      cout << "\tResolution: " << img->getResolution() << endl;
51                      cout << "\tDepth: " << img->getDepth() << endl;
52                      processImageData(img->getRawData(), img->getRawData().size());
53                  }
54              }
55          } catch (Error::Exception &e) {
56              return (EXIT_FAILURE);
57          }
58      }
59      return(EXIT_SUCCESS);
60  }
```

15.2   INCITS Data Records

This INCITS class of data records covers all those record formats that are derived from the standards defined by the InterNational Committee for Information Technology Standards [9]. These formats include the ANSI-2004 Finger Minutiae Record Format [1], the ISO equivalent [2], and other data formats, including finger images.

15.2.1   Finger Views

Within the BECommon, finger view objects (Section 14) can be created from a combination of finger minutiae and image records. However, it is not necessary to have both records in order to create the view because each record contains enough information to represent the finger (image size, for example). However, if a view is contructed using only the minutiae record, then the image itself will not be present. Alternatively, if a view is made from an image record, no minutiae data would be available. It is possible to construct a view without any information.

Listing 13.2 on page 43 shows an example of accessing the information in an ANSI 378-2004 Finger Minutiae Record by creating an ANSI2004View object from the record file.

References

[1] *ANSI INCITS 378-2004: Finger Minutiae Format for Data Interchange*. ANSI/INCITS, 2004. 39, 43, 50

[2] *ISO/IEC 19794-2: Information technology - Biometric data interchange formats - Part 2: Finger minutiae data*. ISO/IEC, first edition, 2005. 39, 50

[3] *American National Standard for Information Systems - Data Format for the Interchange of Fingerprint, Facial, & Other Biometric Information*. ANSI/NIST-ITL, 1-2007 edition, 2007. 4, 39, 41, 47

[4] Mark Adler. zlib, 2012. http://www.zlib.net/. 18

[5] World Wide Web Consortium. Portable Network Graphics Standard, 2003. http://www.w3.org/TR/PNG/. 35

[6] Independent JPEG Group. libjpeg, 2011. http://www.ijg.org/. 34

[7] Joint Photographic Experts Group. JPEG2000 Image Standard, 1992. http://www.jpeg.org/jpeg2000/index.html 34

[8] Joint Photographic Experts Group. JPEG Image Standard, 2011. http://www.jpeg.org/jpeg/index.html 34

[9] International Committee for Information Technology Standards. http://www.incits.org. 43, 50

[10] ISO/IEC Joint Technical Committee 1/SC 37 Biometrics. 43

[11] Communications and Remote Sensing Lab, Université catholique de Louvain. OpenJPEG Library, 2011. http://www.openjpeg.org/. 34

[12] NIST Biometric Image Software, 2011. http://www.nist.gov/itl/iad/ig/nbis.cfm. 7, 34, 35

[13] NIST Image Group. http://www.nist.gov/itl/iad/ig. 1

[14] Greg Roelofs. libpng, 2011. http://www.libpng.org/pub/png/libpng.html. 35

[15] Bjarne Stroustrup. *The C++ Programming Language*. Addison Wesley, special edition, 2000. 3, 9

[16] Wavelet Scalar Quantization Gray-Scale Fingerprint Image Compression Standard, 2010. https://www.fbibiospecs.org/docs/WSQ_Gray-scale_Specification_Version_3_1_Final.pdf. 35